

Reproducing the properties of winning Iterated Prisoner's Dilemma strategies

João Mendes

102079

Instituto Superior Técnico
Portugal

Sara Rijo

102507

Instituto Superior Técnico
Portugal

Sofia Paiva

102835

Instituto Superior Técnico
Portugal

1 Introduction

The Iterated Prisoner's Dilemma (IPD) serves as a foundational model for studying cooperation and competition in multi-agent systems. The article "Properties of Winning Iterated Prisoner's Dilemma Strategies" [1] offers a comprehensive analysis of 195 strategies across diverse tournament settings, identifying key behavioral traits that contribute to success. Reproducing this study is motivated by the need to validate its findings, ensure reproducibility, and explore the robustness of its conclusions under varying conditions. Previous research has often been limited by narrow strategy selections or specific tournament configurations, underscoring the relevance of this extensive evaluation. The primary objectives of this project are:

- **Validation:** Confirm the reproducibility of the original study's results.
- **Extension:** Assess the generalizability of identified winning traits across different environments.
- **Insight:** Provide deeper understanding into the dynamics of strategy performance in IPD scenarios.

2 Approach

The reproduction will utilize the Axelrod-Python library, an open-source framework designed for simulating IPD tournaments. This choice ensures consistency with the original study and leverages a well-supported platform for extensibility.

The visualization of the results will be done using the matplotlib library, which is the most popular for creating many kinds of visualization in Python.

Each strategy is modeled as an autonomous agent capable of making decisions based on historical interactions. The system will simulate tournaments where these agents engage in repeated IPD games, allowing for the observation of emergent behaviors and performance metrics.

System Architecture:

- **Strategy Repository:** A curated collection of 195 strategies, including both deterministic and stochastic variants.
- **Tournament Engine:** A module to configure and execute various tournament types (e.g., standard, noisy, probabilistic ending).
- **Data Analyzer:** Tools to compute performance metrics, such as average score, cooperation rate, and adaptability.
- **Visualization Suite:** Components to generate plots and charts for interpreting results.

Employing the Axelrod-Python library ensures alignment with the original study's methodology, facilitating direct comparisons. The modular architecture allows for scalability and the introduction

of new strategies or tournament conditions, thereby testing the robustness of the original findings.

3 Empirical Evaluation

We will utilize several performance metrics, most being the same or similar to the ones used in the work of Glynatsi et al.

- **Average Score:** Measures the mean payoff per game for each strategy.
- **Cooperation Rate:** Assesses the frequency of cooperative moves.
- **Performance Variance:** Evaluates consistency across different opponents and tournament types.
- **Comparative Cooperation:** Compares the strategy's cooperation with the cooperation of the tournament as a whole.

Validation Process:

- **Replication:** Reproduce the original tournaments and compare results to those reported by Glynatsi et al.
- **Robustness Testing:** Introduce variations in tournament settings (e.g., noise levels, opponent pools) to examine the stability of identified winning traits.
- **Statistical Analysis:** Employ statistical tests to determine the significance of observed differences and validate the reproducibility of key findings.

4 Limitations of the Original Paper

The original study makes two critical assumptions that limit its applicability to real-world multi-agent systems (MAS):

4.1 Fixed Payoff Structure

The paper assumes a static payoff matrix:

$$T = 5, R = 3, P = 1, S = 0 \quad (\text{satisfying } T > R > P > S).$$

While this aligns with canonical IPD literature, it introduces three key constraints:

- **Real-World Abstraction:**
 - In practical MAS, payoffs are rarely quantifiable as fixed numerical values since they are, many times, context-dependent and non-numeric.
- **Sensitivity to Payoff Ratios:**
 - Conclusions may not hold under alternative ratios, e.g.:
 - * $T = 6$: Increased temptation to defect
 - * $R = 4$: Universal cooperation dominance

4.2 Static Noise Model

The study uses homogeneous action-flipping errors but ignores real-world complexity:

- **Dynamic Noise Ignored:**
 - Assumes fixed noise levels (e.g., 5% $C \leftrightarrow D$ flips) throughout tournaments.
 - *However:* Real-world systems face noise that fluctuates over time, requiring adaptive strategies.
- **Heterogeneous Noise Unaddressed:**
 - Uses uniform noise for all agents, ignoring asymmetric errors.
- **Oversimplified Noise Definition:**
 - Noise is reduced to random action flips ($C \leftrightarrow D$), neglecting other real-world perturbations (e.g., delayed actions, payoff uncertainty).

5 Bridging Theory and Practice

To address these limitations, our replication would, ideally:

- (1) Test alternative payoffs.
- (2) Simulate dynamic noise (e.g., oscillating error rates).
- (3) Model heterogeneous noise (e.g., agent-specific flip probabilities).

However, the Axelrod-Python library imposes critical constraints:

- **Fixed Payoffs:** The library imposes the canonical PD payoffs ($T = 5, R = 3, P = 1, S = 0$).
- **Static Noise Model:** Noise is implemented as homogeneous, fixed-probability action flips ($C \leftrightarrow D$) with no option to change for dynamic or heterogeneous perturbations.

To ensure faithful reproduction of the original study, we will:

- Adhere to the library's fixed payoff structure.
- Implement noise strictly as homogeneous action flips ($C \leftrightarrow D$) at fixed rates.

Implications for Critical Analysis:

- These constraints prevent us from fully testing the generalizability of the paper's claims to real-world MAS, where payoffs and noise are context-dependent.
- Our replication will focus on validating the original conclusions within the canonical IPD framework, while explicitly acknowledging this theoretical-reality gap in our discussion.

References

- [1] Nikoleta E. Glynatsi, Vincent Knight, and Marc Harper. 2024. Properties of winning Iterated Prisoner's Dilemma strategies. *PLOS Computational Biology* 20, 12 (12 2024), 1–19. doi:10.1371/journal.pcbi.1012644